

1. Show your model architecture and loss function. (MSE or ...)

```
class AutoEncoder(nn.Module):
    def __init__(self):
        """
        Model initialize
        """
        super(AutoEncoder, self).__init__()
        # Encoder
        self.conv1 = nn.Conv2d(3, 32, 3)
        self.conv2 = nn.Conv2d(32, 64, 3)
        self.pool = nn.MaxPool2d(2, return_indices = True)

        # Decoder
        self.unpool = nn.MaxUnpool2d(2)
        self.deconv2 = nn.ConvTranspose2d(64, 32, 3)
        self.deconv1 = nn.ConvTranspose2d(32, 3, 3)

        # Save for Unpooling
        self.indices = []


        return
```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 32, 24, 24]	896
MaxPool2d-2	[[[-1, 32, 12, 12], [-1, 32, 12, 12]]]	
Conv2d-3	[-1, 64, 10, 10]	18,496
MaxPool2d-4	[[[-1, 64, 5, 5], [-1, 64, 5, 5]]]	
MaxUnpool2d-5	[-1, 64, 10, 10]	0
ConvTranspose2d-6	[-1, 32, 12, 12]	18,464
MaxUnpool2d-7	[-1, 32, 24, 24]	0
ConvTranspose2d-8	[-1, 3, 26, 26]	867
=====		
Total params:	38,723	
Trainable params:	38,723	
Non-trainable params:	0	

Input size (MB):	0.01	
Forward/backward pass size (MB):	181.10	
Params size (MB):	0.15	
Estimated Total Size (MB):	181.26	

Optimizing for MSE means your generated output intensities are *symmetrically close* to the input intensities. A higher-than-training intensity is penalized by the same amount as an equally valued lower intensity.

2. Plot training loss.



3. Visualize 5 generated samples for each class like this:

